

IT5100A

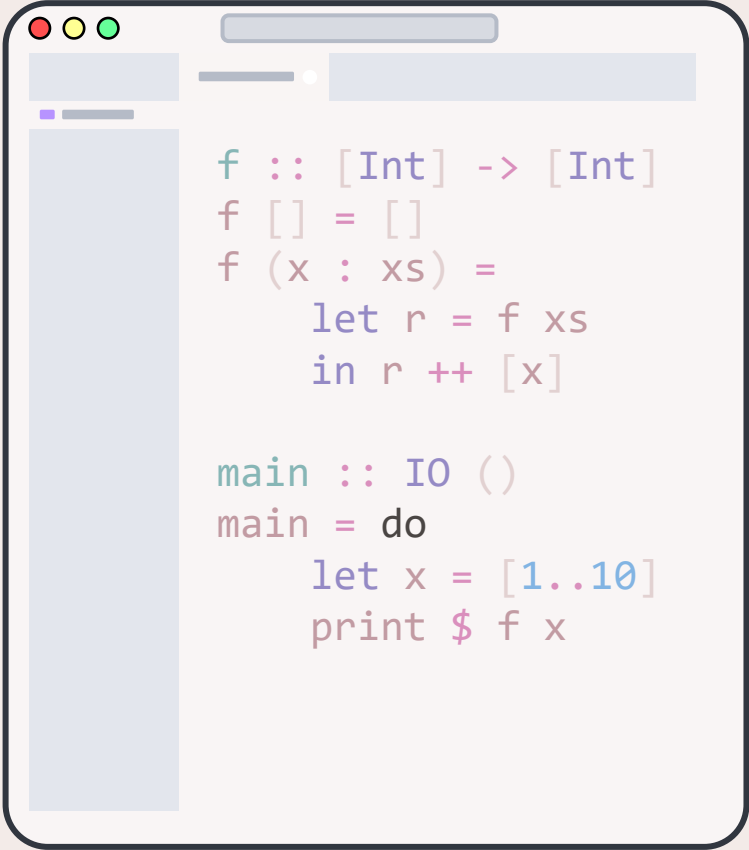
Industry Readiness:

Typed Functional Programming

Concurrent & Parallel Programming

Foo Yong Qi


yongqi@nus.edu.sg



```
f :: [Int] -> [Int]
f [] = []
f (x : xs) =
    let r = f xs
    in r ++ [x]

main :: IO ()
main = do
    let x = [1..10]
    print $ f x
```

Concurrent Programming



Multiple tasks done “at the same time”


Parallel Programming



One task split into smaller tasks and done “at the same time”

Concurrent Programming

Concurrent Programming



Each task can be provided as a single **thread** of control

- Use `forkIO` to create new threads of control
- Threads are nondeterministic

```
ghci> import Control.Concurrent
ghci> :t forkIO
forkIO :: IO () -> IO ThreadId
ghci> import System.Directory
ghci> forkIO (writeFile "it5100a-notes.md" "Hello World!")
      >> doesFileExist "it5100a-notes.md"
False
```

Concurrent programs “hide latency”

```
import Data.List

writeContents :: String -> String -> IO ()
writeContents file_name contents = do
    let c = intercalate "\n" $
        replicate 1000000
            (intercalate "" $ replicate 100 contents)
    writeFile file_name c
    putStrLn $ file_name ++ " written"
```

Concurrent programs “hide latency”

```
main :: IO ()
main = do
  putStrLn "Enter filename:"
  x <- getLine
  if x == "exit"
  then return ()
  else do putStrLn "Enter file contents:"
          y <- getLine
          forkIO $ writeContents x y
          main
```

MVars: synchronizing mutable variables

```
ghci> import Control.Concurrent.MVar
ghci> :t putMVar
putMVar :: MVar a -> a -> IO ()
ghci> :t takeMVar
takeMVar :: MVar a -> IO a
```


Can use **MVar** for obtaining/storing information shared between threads

```
writeContents :: String -> String -> MVar () -> IO ()
writeContents file_name contents lock = do
    takeMVar lock
    putStrLn $ "Write to " ++ file_name ++ " started"
    let !c = intercalate "\n" $
        replicate 1000000 (intercalate "" $
            replicate 500 contents)
    writeFile file_name c
    putStrLn $ file_name ++ " written"
    putMVar lock ()
```

- Use the **MVar** as a lock, only exit when **all** writing threads are done
- Haskell runtime wakes up sleeping threads in FIFO order

```
mainLoop :: MVar () -> IO ()
mainLoop lock = do
    putStrLn "Enter filename:"
    x <- getLine
    if x == "exit"
    then do takeMVar lock
            return ()
    else do putStrLn "Enter file contents:"
            y <- getLine
            forkIO $ writeContents x y lock
            mainLoop lock
```

Chans: one-way communication channels

```
ghci> import Control.Concurrent.Chan
ghci> :t writeChan
writeChan :: Chan a -> a -> IO ()
ghci> :t readChan
readChan :: Chan a -> IO a
```

Can use **Chan** to pass information from main thread to writing thread

```
writeContents :: Chan String -> IO ()
writeContents chan = do
    file_name <- readChan chan
    contents   <- readChan chan
    putStrLn $ "Write to " ++ file_name ++ " started"
    let c = intercalate "\n" $
        replicate 1000000 (
            intercalate "" $ replicate 500 contents)
    writeFile file_name c
    putStrLn $ file_name ++ " written"
    writeContents chan
```

Can use **Chan** to pass information from main thread to writing thread

```
mainLoop :: Chan String -> IO ()
mainLoop chan = do
    putStrLn "Enter filename:"
    x <- getLine
    if x == "exit"
    then return ()
    else do putStrLn "Enter file contents:"
            y <- getLine
            writeChan chan x
            writeChan chan y
    mainLoop chan
```

Concurrency doesn't need multiple cores, but can certainly help!
For Haskell, include **threaded runtime**

```
ghc Main.hs -threaded  
./Main  
4
```

```
import Control.Concurrent  
main :: IO ()  
main = do setNumCapabilities 4  
          print numCapabilities
```

Parallel Programming

Parallel Programming



Can be done with concurrency features, but Haskell is unique due to **non-strict evaluation**

- When expressions are defined, Haskell puts a thunk
- Thunks are evaluated **on-demand** up to **Head Normal Form (HNF)**

```
ghci> x = [1..]  
ghci> case x of { [] -> 0; (x:xs) -> x }  
1  
ghci> let x = [1..]; y = sum x in 1 + 2  
3
```

Non-strict evaluation means some times we accidentally run computations on the main thread, not on forked threads

```
expensive :: MVar String -> IO ()
expensive var = do
    putMVar var expensivelyComputedString

main :: IO ()
main = do
    var <- newEmptyMVar
    forkIO $ expensive var
    whatever
    result <- takeMVar var
    print result
```

`seq` introduces an artificial demand on first argument for evaluation
on second argument

```
ghci> :t seq
seq :: a -> b -> b
```

```
ghci> let x = [1..]; y = sum x in y `seq` 1 + 2
-- ... does not terminate
ghci> let x = [1..] in x `seq` 1 + 2
3 -- ... only evaluates x to HNF
```

Example: create new evaluation strategy that evaluates deeply

```
ghci> :{
ghci| deepSeq :: [a] -> b -> b
ghci| deepSeq [] x = x
ghci| deepSeq (x:xs) y = x `seq` deepSeq xs y
ghci| :}
ghci> x = [1..]
ghci> x `seq` 1
1
ghci> x `deepSeq` 1
-- does not terminate
```

Parallel programming can be done by using **parallel evaluation strategies**

```
ghci> import GHC.Conc
ghci> :t par
par :: a -> b -> b
ghci> :t pseq
pseq :: a -> b -> b
```

$x \text{ `par` } (f \ x \ y)$

Evaluate x in parallel with $f \ x \ y$; what if evaluation of $f \ x \ y$ happens first which evaluates x first? No parallelism happens

```
x `par` (y `pseq` f x y)
```

Evaluate **x** in parallel, evaluate **y** before **f x y** on the main thread

Parallel fibonacci

```
fib :: Int -> Integer
fib 0 = 0
fib 1 = 1
fib n = n1 `par` (n2 `pseq` (n1 + n2))
    where n1 = fib (n - 1)
          n2 = fib (n - 2)
```



```
> time cabal run playground -- +RTS -N20
```

```
Number of cores: 20
```

```
1134903170
```

Executed in	3.29 secs	fish	external
usr time	53.14 secs	319.00 micros	53.14 secs
sys time	0.47 secs	129.00 micros	0.47 secs

```
> time cabal run playground -- +RTS -N1
```

```
Number of cores: 1
```

```
1134903170
```

Executed in	12.93 secs	fish	external
usr time	12.61 secs	418.00 micros	12.61 secs
sys time	0.08 secs	171.00 micros	0.08 secs

At some point, parallel evaluation overhead outweighs performance gains, put a threshold

```
fib :: Int -> Integer
fib 0 = 0
fib 1 = 1
-- sequential for small n
fib n | n <= 10 = fib (n - 1) + fib (n - 2)
-- parallel for large n
fib n = n1 `par` (n2 `pseq` (n1 + n2))
  where n1 = fib (n - 1)
        n2 = fib (n - 2)
```

```
> time cabal run playground -- +RTS -N20
```

```
Number of cores: 20
```

```
1134903170
```

Executed in	892.37 millis	fish	external
usr time	13.01 secs	646.00 micros	13.00 secs
sys time	0.18 secs	0.00 micros	0.18 secs

```
> time cabal run playground -- +RTS -N1
```

```
Number of cores: 1
```

```
1134903170
```

Executed in	6.81 secs	fish	external
usr time	6.71 secs	453.00 micros	6.71 secs
sys time	0.03 secs	0.00 micros	0.03 secs

Parallel merge sort

```
mergesort :: Ord a => [a] -> [a]
mergesort [] = []
mergesort [x] = [x]
mergesort ls
  | n < 100 = merge left' right'
  | otherwise = par left' $ pseq right' $ merge left' right'
  where n = length ls `div` 2
        merge [] ys = ys
        merge xs [] = xs
        merge (x:xs) (y:ys)
          | x <= y = x : merge xs (y : ys)
          | otherwise = y : merge (x:xs) ys
        (left, right) = splitAt n ls
        left' = mergesort left
        right' = mergesort right
```

```
> time cabal run playground -- +RTS -N20
```

```
Number of cores: 20
```

```
10000000
```

Executed in	3.58 secs	fish	external
usr time	16.02 secs	381.00 micros	16.02 secs
sys time	1.39 secs	159.00 micros	1.39 secs

```
> time cabal run playground -- +RTS -N1
```

```
Number of cores: 1
```

```
10000000
```

Executed in	6.11 secs	fish	external
usr time	5.62 secs	0.00 micros	5.62 secs
sys time	0.43 secs	586.00 micros	0.43 secs

Force full evaluation in other thread

```
mergesort :: Ord a => [a] -> [a]
mergesort [] = []
mergesort [x] = [x]
mergesort ls
  | n < 100 = merge left' right'
  | otherwise = par (deepSeq left') $ pseq right' $ merge left' right'
  where n = length ls `div` 2
        merge [] ys = ys
        merge xs [] = xs
        merge (x:xs) (y:ys)
          | x <= y = x : merge xs (y : ys)
          | otherwise = y : merge (x:xs) ys
        (left, right) = splitAt n ls
        left' = mergesort left
        right' = mergesort right

deepSeq :: [a] -> ()
deepSeq [] = ()
deepSeq (x:xs) = x `seq` deepSeq xs
```

```
> time cabal run playground -- +RTS -N20
```

```
Number of cores: 20
```

```
10000000
```

Executed in	2.89 secs	fish	external
usr time	18.04 secs	365.00 micros	18.04 secs
sys time	0.68 secs	145.00 micros	0.67 secs

```
> time cabal run playground -- +RTS -N1
```

```
Number of cores: 1
```

```
10000000
```

Executed in	6.18 secs	fish	external
usr time	5.59 secs	362.00 micros	5.59 secs
sys time	0.46 secs	145.00 micros	0.46 secs

Can use `parallel` library to express parallel evaluation strategies easily

```
otherwise = runEval $ do
  l <- rparWith rdeepseq left'
  r <- rseq right'
  return $ merge l r
```

```
fib n = runEval $ do
  n1 <- rpar (fib (n - 1))
  n2 <- rseq (fib (n - 2))
  return $ n1 + n2
```


Separate algorithm from evaluation strategy

```
fib :: Int -> Integer
fib 0 = 0
fib 1 = 1
fib n | n <= 10    = n1 + n2
      | otherwise = (n1 + n2) `using` strat
  where n1 = fib (n - 1)
        n2 = fib (n - 2)
        strat v = do { rpar n1; rseq n2; return v }
```

Software Transactional Memory

What problems can happen?

```
swap a b chan = do
  x <- takeMVar a
  y <- takeMVar b
  putMVar a y
  putMVar b x
  writeChan chan ()
```

```
addToMVar a b chan = do
  y <- takeMVar b
  x <- takeMVar a
  let z = x + y
  putMVar b z
  putMVar a x
  writeChan chan ()
```

Example transaction: taking two mutable variables; STM equivalent of MVar is TMVar

```
takeBothTMVars :: TMVar a -> TMVar b -> STM (a, b)
takeBothTMVars a b = do
  x <- takeTMVar a
  y <- takeTMVar b
  return (x, y)
```

```
takeBothMVars :: MVar a -> MVar b -> IO (a, b)
takeBothMVars a b = do
  x <- takeMVar a
  y <- takeMVar b
  return (x, y)
```

Both **TMVars** are taken in one fell swoop, no lock order inversion since atomic transaction is indivisible

```
swap :: TMVar a -> TMVar a -> Chan () -> IO ()
swap a b chan = do
    (x, y) <- atomically $ takeBothTMVars a b
    -- ...

addToMVar :: Num a => TMVar a -> TMVar a
           -> Chan () -> IO ()
addToMVar a b chan = do
    (y, x) <- atomically $ takeBothTMVars b a
    -- ...
```

- Concurrency not very different in Haskell than other languages, yet not so applicable
- Parallelism easy due to non-strict evaluation in Haskell
- Haskell concurrency has same problems with other languages
- STM monad makes composing transactions easy

Thank you

Foo Yong Qi

yongqi@nus.edu.sg

<https://yongqi.foo/>

